


Introduction to Neural Networks

LAB # 11 & 12



Fall 2021

Data Analytics Lab

Submitted by: **Shah Raza**

Registration No.: **18PWCSE1658**

Class Section: **B**

“On my honor, as student of University of Engineering and Technology, I have neither given nor received unauthorized assistance on this academic work.”

Student Signature: _____

Submitted to:

Engr. Mian Ibad Ali Shah

9 March 2022

Department of Computer Systems Engineering

University of Engineering and Technology, Peshawar

Introduction to Neural Network:


Before that, let's study a little mathematics (Linear Algebra)

Scalar (0D)

Vector (1D)


Matrix (2D) -> combination of vectors

Tensor -> collection of matrices

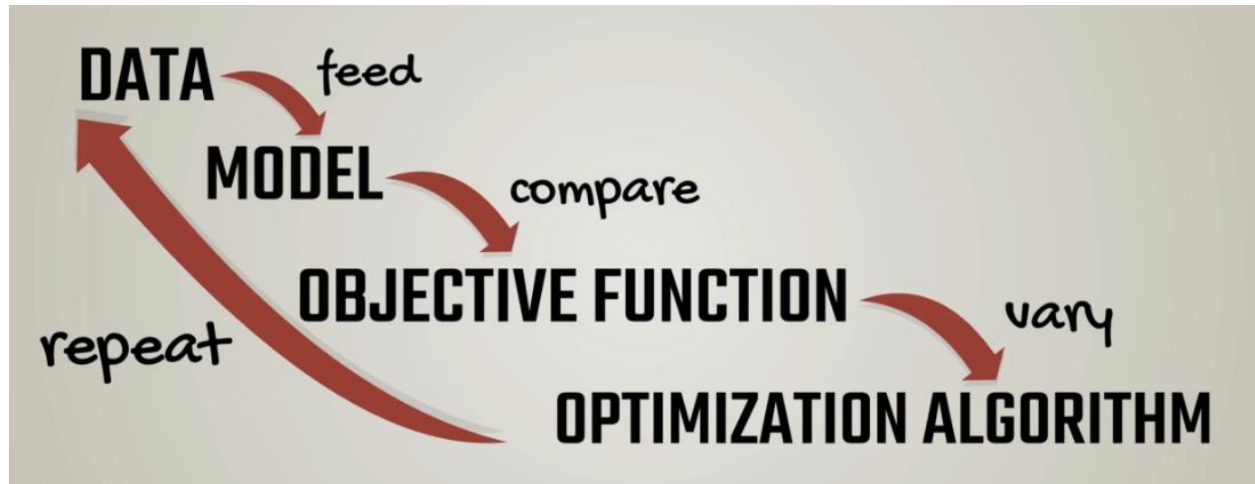


Transposing an array (vector, matrix or tensor) -> `x.T` in python

Dot product (scalar product) -> `np.dot(x,y)`



Ingredients of an Algorithm:



Training the model:

Coffee machine making coffee

Self-driving cars

Types of Machine Learning:

Supervised

Unsupervised

Reinforcement Learning

What is a model?

Example: Linear Regression

$Y=xw+b$

Objective Function:

Objective function is the measure used to evaluate how well the model's outputs match the desired correct values.

It has two types:

- Loss Function (Example: Supervised Learning) aka cost function
 - For example: Regression cost function: L2 norm.
 - Classification: Cross entropy
- Reward Function (Example: Reinforcement Learning)

Optimization Algorithm: Example, Gradient Descent

Task 1:

Choosing the objective function and the optimization method

```
In [27]: # Again, we use a loss function. mean_squared_error is the scaled L2-norm (per observation)
mean_loss = tf.losses.mean_squared_error(labels=targets, predictions=outputs) / 2.
# Note that there also exists a function tf.nn.L2_loss.

# Instead of implementing Gradient Descent on our own, in TensorFlow we can simply state
# "Minimize the mean loss by using Gradient Descent with a given learning rate"
# Simple!
optimize = tf.train.GradientDescentOptimizer(learning_rate=0.05).minimize(mean_loss)
```

Prepare for execution

```
In [28]: # So far we've defined the placeholders, variables, the loss function and the optimization method.
# The actual training happens inside sessions.
sess = tf.InteractiveSession()
```

Initializing variables

```
In [29]: # Before we start training, we need to initialize our variables: the weights and biases.
initializer = tf.global_variables_initializer()

sess.run(initializer)
```

Activate Windows
Go to Settings to activate

Learning

```
In [31]: # As in the previous lab, we train for a set number (100) of iterations over the dataset
for i in range(100):
    # sess.run is the session's function to actually do something, anything.
    # Above, we used it to initialize the variables.
    # Here, we use it to feed the training data to the computational graph, defined by the feed_dict parameter
    # So the line of code means: "Run the optimize and mean_loss operations by filling the placeholder
    # objects with data from the feed_dict parameter".
    _, curr_loss = sess.run([optimize, mean_loss],
                           feed_dict={inputs: training_data['inputs'], targets: training_data['targets']})


    print(curr_loss)
```

```
224.36153
96.79172
45.56499
24.258543
14.921057
10.505043
8.188638
6.8132415
5.8881536
5.197546
4.6420927
4.1735177
3.7667549
3.4077315
3.087803
2.801146
2.5434911
2.311103
```

Activate Windows
Go to Settings to activate W

Plotting the data

```
In [32]: out = sess.run([outputs],
                        feed_dict={inputs: training_data['inputs']})
plt.plot(np.squeeze(out), np.squeeze(training_data['targets']))
plt.xlabel('outputs')
plt.ylabel('targets')
plt.show()
```



In []:

Activate Windows
Go to Settings to activate W

Task 2:

Simple Linear Regression. Minimal example

Import the relevant libraries

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits.mplot3d import Axes3D
```

Generate random input data to train on

```
In [2]: observations = 1000
xs = np.random.uniform(low=-10, high=10, size=(observations,1))
zs = np.random.uniform(-10, 10, (observations,1))

inputs = np.column_stack((xs,zs))

print (inputs.shape)

(1000, 2)
```

Generate the targets we will aim at

```
In [3]: # We want to make a function and see if the algorithm has learned it.
# We add a small random noise to the function i.e.  $f(x,z) = 2x - 3z + 5$  + <small noise> as in real life datasets
noise = np.random.uniform(-1, 1, (observations,1))
```

Plot the training data

The point is to see that there is a strong trend that our model should learn to reproduce.


```
In [7]: targets = targets.reshape(observations,)

fig = plt.figure()

ax = fig.add_subplot(111, projection='3d')

ax.plot(xs, zs, targets)

ax.set_xlabel('xs')
ax.set_ylabel('zs')
ax.set_zlabel('Targets')
ax.view_init(azim=250)
plt.show()
targets = targets.reshape(observations,1)
```



Set a learning rate

```
In [9]: # Play around with learning rate.
learning_rate = 0.05
```

Train the model

```
In [14]: for i in range(1000):

    outputs = np.dot(inputs, weights) + biases
    deltas = outputs - targets

    loss = np.sum(deltas ** 2) / 2 / observations

    print (loss)

    deltas_scaled = deltas / observations

    # Finally, apply the gradient descent update rules
    # The weights are 2x1, learning rate is 1x1 (scalar), inputs are 1000x2, and deltas_scaled are 1000x1
    # Transpose the inputs so that we get an allowed operation.
    weights = weights - learning_rate * np.dot(inputs.T, deltas_scaled)
    biases = biases - learning_rate * np.sum(deltas_scaled)
```

```
0.1682932421214414
0.16813935858891543
0.1679915681123002
0.16784962943556467
0.16771331085526878
.....
```

Activate Windows
Go to Settings to activate


```
In [15]: print (weights, biases)

[[ 2.00372083]
 [-3.00100363]] [5.00604098]
```

Plot last outputs vs targets

Since they are the last ones at the end of the training, they represent the final model accuracy.
The closer this plot is to a 45 degree line, the closer target and output values are.

```
In [16]: plt.plot(outputs, targets)
plt.xlabel('outputs')
plt.ylabel('targets')
plt.show()
```



Activate Windows
Go to Settings to activate

